

Investigate_a_Dataset

December 31, 2022

1

1.1 Introduction

1.1.1 For my project, I am going to analyse the TMDb movie data set. This dataset contains about 10,000 movies collected from The Movie Database (TMDb). A quick scan of the dataset csv file can show us that it has about 10,000 rows and 21 columns. I can also notice that some columns have null values for example release_date column.

2 Questions

2.0.1 Some questions I asked myself are, "Which year produced the highest release of movies?" and "Which movie genre is the produced the highest and lowest number of movies?" These questions are going to help me explore the data in depth.

```
In [28]: # Use this cell to set up import statements for all of the packages that you
        #      plan to use.
        import pandas as pd
        import numpy as np
        import matplotlib.pyplot as plt
        import seaborn as sns
        % matplotlib inline

        # Remember to include a 'magic word' so that your visualizations are plotted
        #      inline with the notebook. See this page for more:
        #      http://ipython.readthedocs.io/en/stable/interactive/magics.html

In [29]: # Upgrade pandas to use dataframe.explode() function.
        !pip install --upgrade pandas==0.25.0
```

```
Requirement already up-to-date: pandas==0.25.0 in /opt/conda/lib/python3.6/site-packages (0.25.0)
Requirement already satisfied, skipping upgrade: pytz>=2017.2 in /opt/conda/lib/python3.6/site-packages
Requirement already satisfied, skipping upgrade: numpy>=1.13.3 in /opt/conda/lib/python3.6/site-packages
Requirement already satisfied, skipping upgrade: python-dateutil>=2.6.1 in /opt/conda/lib/python3.6/site-packages
Requirement already satisfied, skipping upgrade: six>=1.5 in /opt/conda/lib/python3.6/site-packages
```

Data Wrangling

3 In this section of the report, you will load in the data, check for cleanliness, and then trim and clean your dataset for analysis.

```
In [30]: # Load your data and print out a few lines. Perform operations to inspect data
df = pd.read_csv('tmdb-movies.csv')
df.head(5)
```

```
Out[30]:
```

	id	imdb_id	popularity	budget	revenue	\
0	135397	tt0369610	32.985763	150000000	1513528810	
1	76341	tt1392190	28.419936	150000000	378436354	
2	262500	tt2908446	13.112507	110000000	295238201	
3	140607	tt2488496	11.173104	200000000	2068178225	
4	168259	tt2820852	9.335014	190000000	1506249360	

	original_title	\
0	Jurassic World	
1	Mad Max: Fury Road	
2	Insurgent	
3	Star Wars: The Force Awakens	
4	Furious 7	

	cast	\
0	Chris Pratt Bryce Dallas Howard Irrfan Khan Vi...	
1	Tom Hardy Charlize Theron Hugh Keays-Byrne Nic...	
2	Shailene Woodley Theo James Kate Winslet Ansel...	
3	Harrison Ford Mark Hamill Carrie Fisher Adam D...	
4	Vin Diesel Paul Walker Jason Statham Michelle ...	

	homepage	director	\
0	http://www.jurassicworld.com/	Colin Trevorrow	
1	http://www.madmaxmovie.com/	George Miller	
2	http://www.thedivergentseries.movie/#insurgent	Robert Schwentke	
3	http://www.starwars.com/films/star-wars-episod...	J.J. Abrams	
4	http://www.furious7.com/	James Wan	

	tagline	...	\
0	The park is open.	...	
1	What a Lovely Day.	...	
2	One Choice Can Destroy You	...	
3	Every generation has a story.	...	
4	Vengeance Hits Home	...	

	overview	runtime	\
0	Twenty-two years after the events of Jurassic ...	124	
1	An apocalyptic story set in the furthest reach...	120	
2	Beatrice Prior must confront her inner demons ...	119	
3	Thirty years after defeating the Galactic Empi...	136	
4	Deckard Shaw seeks revenge against Dominic Tor...	137	

```

                                genres \
0  Action|Adventure|Science Fiction|Thriller
1  Action|Adventure|Science Fiction|Thriller
2      Adventure|Science Fiction|Thriller
3  Action|Adventure|Science Fiction|Fantasy
4      Action|Crime|Thriller

```

```

                                production_companies release_date vote_count \
0  Universal Studios|Amblin Entertainment|Legenda...      6/9/15      5562
1  Village Roadshow Pictures|Kennedy Miller Produ...      5/13/15      6185
2  Summit Entertainment|Mandeville Films|Red Wago...      3/18/15      2480
3      Lucasfilm|Truenorth Productions|Bad Robot      12/15/15      5292
4  Universal Pictures|Original Film|Media Rights ...      4/1/15      2947

```

```

      vote_average  release_year  budget_adj  revenue_adj
0              6.5           2015  1.379999e+08  1.392446e+09
1              7.1           2015  1.379999e+08  3.481613e+08
2              6.3           2015  1.012000e+08  2.716190e+08
3              7.5           2015  1.839999e+08  1.902723e+09
4              7.3           2015  1.747999e+08  1.385749e+09

```

[5 rows x 21 columns]

```
In [31]: # types and look for instances of missing or possibly errant data
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10866 entries, 0 to 10865
Data columns (total 21 columns):
id                10866 non-null int64
imdb_id           10856 non-null object
popularity        10866 non-null float64
budget            10866 non-null int64
revenue           10866 non-null int64
original_title    10866 non-null object
cast              10790 non-null object
homepage          2936 non-null object
director          10822 non-null object
tagline           8042 non-null object
keywords          9373 non-null object
overview          10862 non-null object
runtime           10866 non-null int64
genres            10843 non-null object
production_companies 9836 non-null object
release_date      10866 non-null object
vote_count        10866 non-null int64
vote_average      10866 non-null float64

```

```
release_year          10866 non-null int64
budget_adj            10866 non-null float64
revenue_adj           10866 non-null float64
dtypes: float64(4), int64(6), object(11)
memory usage: 1.7+ MB
```

```
In [32]: #show the rows and columns of the dataset
df.shape
```

```
Out[32]: (10866, 21)
```

```
In [33]: # values of missing data per column
df.isnull().sum()
```

```
Out[33]: id                0
imdb_id                  10
popularity               0
budget                  0
revenue                 0
original_title          0
cast                   76
homepage                7930
director                44
tagline                 2824
keywords               1493
overview                4
runtime                 0
genres                  23
production_companies    1030
release_date            0
vote_count              0
vote_average            0
release_year            0
budget_adj              0
revenue_adj             0
dtype: int64
```

```
In [34]: df.duplicated().sum()
```

```
Out[34]: 1
```

4 Data Cleaning

4.1 Here we remove information that we don't need to use during our analysis

```
In [35]: # After discussing the structure of the data and any problems that need to be
# cleaned, perform those cleaning steps in the second part of this section.
#drop the duplicated rows
```

```
df.drop_duplicates(inplace=True)
df.shape
```

```
Out[35]: (10865, 21)
```

```
In [36]: #drop the columns that I will not use in my analysis
df.drop(columns = ['imdb_id', 'homepage', 'production_companies', 'tagline', 'keywords'])
df.shape
```

```
Out[36]: (10865, 16)
```

```
In [37]: #check missing values
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 10865 entries, 0 to 10865
Data columns (total 16 columns):
id                10865 non-null int64
popularity        10865 non-null float64
budget            10865 non-null int64
revenue           10865 non-null int64
original_title    10865 non-null object
cast              10789 non-null object
director          10821 non-null object
overview          10861 non-null object
runtime           10865 non-null int64
genres            10842 non-null object
release_date      10865 non-null object
vote_count        10865 non-null int64
vote_average      10865 non-null float64
release_year      10865 non-null int64
budget_adj        10865 non-null float64
revenue_adj       10865 non-null float64
dtypes: float64(4), int64(6), object(6)
memory usage: 1.4+ MB
```

```
In [38]: #handle missing values by filling in missing values with the previous valid value in the column
df.fillna(method='ffill', inplace=True)
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 10865 entries, 0 to 10865
Data columns (total 16 columns):
id                10865 non-null int64
popularity        10865 non-null float64
budget            10865 non-null int64
revenue           10865 non-null int64
original_title    10865 non-null object
```

```

cast                10865 non-null object
director            10865 non-null object
overview            10865 non-null object
runtime             10865 non-null int64
genres              10865 non-null object
release_date        10865 non-null object
vote_count          10865 non-null int64
vote_average        10865 non-null float64
release_year        10865 non-null int64
budget_adj          10865 non-null float64
revenue_adj         10865 non-null float64
dtypes: float64(4), int64(6), object(6)
memory usage: 1.4+ MB

```

```

In [39]: #change date format
         df['release_date'] = pd.to_datetime(df['release_date'])
         df['release_date'].head()

```

```

Out[39]: 0    2015-06-09
         1    2015-05-13
         2    2015-03-18
         3    2015-12-15
         4    2015-04-01
         Name: release_date, dtype: datetime64[ns]

```

```

In [40]: df.info()

```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 10865 entries, 0 to 10865
Data columns (total 16 columns):
id                10865 non-null int64
popularity        10865 non-null float64
budget            10865 non-null int64
revenue           10865 non-null int64
original_title    10865 non-null object
cast              10865 non-null object
director          10865 non-null object
overview          10865 non-null object
runtime           10865 non-null int64
genres            10865 non-null object
release_date      10865 non-null datetime64[ns]
vote_count        10865 non-null int64
vote_average      10865 non-null float64
release_year      10865 non-null int64
budget_adj        10865 non-null float64
revenue_adj       10865 non-null float64
dtypes: datetime64[ns](1), float64(4), int64(6), object(5)

```

memory usage: 1.4+ MB

Exploratory Data Analysis

4.2 Here we Compute statistics and create visualizations with the goal of addressing the research questions at the Introduction section.


4.2.1 Research Question 1 - Which year produced the highest release of movies?

```
In [41]: # Use this, and more code cells, to explore your data. Don't forget to add
#         Markdown cells to document your observations and findings.
#group data by release-year column then find the totals in terms of the id
df_movies = df.groupby('release_year').count()['id']
df_movies.tail()
```

```
Out[41]: release_year
2011      540
2012      588
2013      659
2014      700
2015      629
Name: id, dtype: int64
```

```
In [42]: #plot the data for viusalisation
df_movies.plot(xticks = np.arange(1960,2015,5), figsize = (14,5))
plt.title("Number of movies released per year", fontsize = 10)
plt.xlabel("Release Year", fontsize = 10)
plt.ylabel("Number of Movies", fontsize = 10)
```

```
Out[42]: Text(0,0.5,'Number of Movies')
```



4.2.2 Research Question 2 - Which movie genre is the produced the highest and lowest number of movies?

```
In [43]: # Continue to explore the data to address your additional research
# questions. Add more headers as needed if you have more questions to
# investigate.
# some entries in the genre column contain multiple entries. We will split the out the
df_genre = df
# columns to split by "/"
df_genre['genres'] = df['genres'].apply(lambda x: x.split("/")[0])
```

```
In [44]: # Confirm for split
df_genre.genres.unique()
```

```
Out[44]: array(['Action', 'Adventure', 'Western', 'Science Fiction', 'Drama',
               'Family', 'Comedy', 'Crime', 'Romance', 'War', 'Mystery',
               'Thriller', 'Fantasy', 'History', 'Animation', 'Horror', 'Music',
               'Documentary', 'TV Movie', 'Foreign'], dtype=object)
```

```
In [45]: #we use groupby to group the genre column then find total number of each movie and also
df_genre = df.groupby(['genres']).count()['id'].sort_values(ascending=False)
df_genre
```

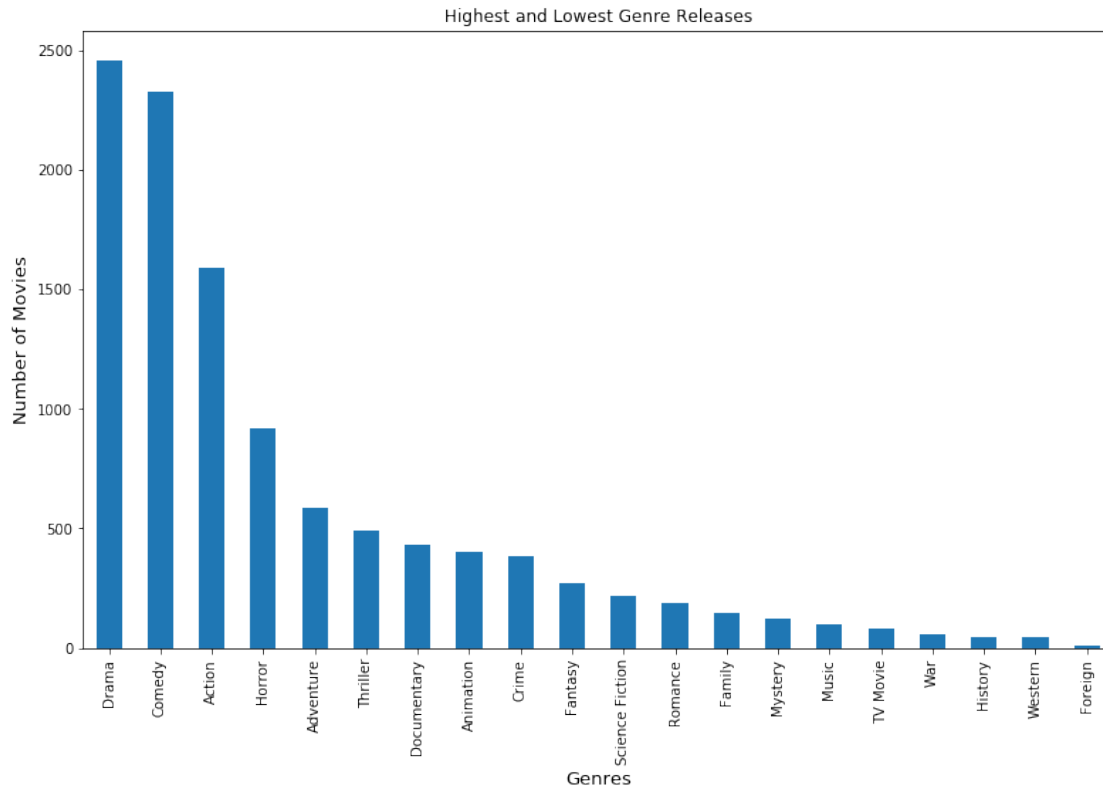
```
Out[45]: genres
Drama                2459
Comedy               2324
Action              1591
Horror               916
Adventure            586
Thriller             492
Documentary          433
Animation            404
Crime                382
Fantasy              272
Science Fiction      216
Romance              186
Family               145
Mystery              125
Music                100
TV Movie              79
War                   59
History               44
Western               43
Foreign               9
Name: id, dtype: int64
```

```
In [49]: #visualization
df_genre.plot(kind = 'bar', figsize = (13,8))
plt.title('Highest and Lowest Genre Releases', fontsize = 12)
```



```
plt.xlabel('Genres', fontsize = 13)
plt.ylabel('Number of Movies', fontsize = 13)
```

```
Out[49]: Text(0,0.5,'Number of Movies')
```



Conclusions

- 5 After careful analysis of this dataset, we can finally answer our questions. The year 2014 produced the highest release of movies. In addition to that, we can also see that the graph is skewed to the right meaning the release of movies increased as the years went by.
- 6 Drama Genre produced the highest number of movies while Foreign goes in with the least number of movies produced.
- 7 Limitations
- 8 During cleaning, We filled up some of the missing data in columns such as the genre column. This might affect the validity of the results we accrued.

```
In [52]: from subprocess import call  
         call(['python', '-m', 'nbconvert', 'Investigate_a_Dataset.ipynb'])
```

```
Out[52]: 0
```

```
In [ ]:
```